

# Jacobi method Implementation in Spark


---

Azam Asl

September 28, 2017

Computer Science Department

## Parallel computing!



# Project overview

- We implemented Jacobi method to solve a Laplace equation in Spark + YARN in three different implementation.

# Project overview

- We implemented Jacobi method to solve a Laplace equation in Spark + YARN in three different implementation.
- First implementation is a naive implementation using dense matrix. It didn't scale for large discretizations ( $N \geq 10^5$ ).

# Project overview

- We implemented Jacobi method to solve a Laplace equation in Spark + YARN in three different implementation.
- First implementation is a naive implementation using dense matrix. It didn't scale for large discretizations ( $N \geq 10^5$ ).
- The second implementation which uses sparse matrix scales at least 10 times  $N$  in compare to our first implementation.

# Project overview

- We implemented Jacobi method to solve a Laplace equation in Spark + YARN in three different implementation.
- First implementation is a naive implementation using dense matrix. It didn't scale for large discretizations ( $N \geq 10^5$ ).
- The second implementation which uses sparse matrix scales at least 10 times  $N$  in compare to our first implementation.
- The third implementation is an improvement on second version and it runs faster on my laptop.
- Testing on YARN cluster for  $N > 10^6$  was useless. Supposedly because of unresolved issue that Spark application hangs when dynamic allocation is enabled in Hadoop enviroment.

# Outline

- Problem statement.
- Implementation details.
- Spark UI demo.

# Problem statement: 1D Laplace equation (aka. boundary value problem)

Given a function  $f : [0, 1] \rightarrow \mathbb{R}$  we would like to find function  $u$  such that its second derivative within the interval  $(0, 1)$  is equal to  $-f$  and its boundary values are 0. More precisely:

$$u'' = -f \text{ in } (0, 1) \text{ and } u(0) = 0, u(1) = 0$$

Using finite-difference approximations for the second derivative to approximate the solution, Using Taylor expansions of  $u(x_i - h)$  and  $u(x_i + h)$  where :

$\{x_i = ih : i = 0, 1, \dots, N, N+1\} \subset [0, 1]$ , with  $h = \frac{1}{N+1}$ :

$$\frac{1}{h^2} \begin{bmatrix} 2 & -1 & 0 & \dots & 0 \\ -1 & 2 & -1 & \dots & \vdots \\ 0 & \ddots & \ddots & \ddots & 0 \\ \vdots & & -1 & 2 & -1 \\ 0 & \dots & 0 & -1 & 2 \end{bmatrix} \begin{bmatrix} u_1 \\ u_2 \\ \vdots \\ u_{N-1} \\ u_N \end{bmatrix} = \begin{bmatrix} f_1 \\ f_2 \\ \vdots \\ f_{N-1} \\ f_N \end{bmatrix}$$



## Solving Laplace equation with Jacobi

In every iteration computes  $u^{k+1}$  (i.e. update) from  $u^k$  per below (element-wise):

$$u_i^{k+1} = \frac{1}{a_{ii}}(f_i - \sum_{j \neq i} a_{ij} u_j^k)$$

**Data:** compute  $hsq \leftarrow h$  square, set  $max\_iter$ , set  $tol$

**Result:** updated  $u$  and residual

- 1 Initialization square matrix  $A$ , column vector  $u$ , set maximum number of iterations;
- 2 **while**  $cur\_iter \leq max\_iter$  and  $residual \geq tol$  **do**
- 3   Compute vector  $v_1 = Au$ ;
- 4   Update each element of  $u = u + hsq * 0.51 - 0.5v_1$  ;
- 5   Compute  $v_2 = Au$  ;
- 6   Set residual to L2 norm of  $(\frac{1}{hsq}v_2 - f)$ ;
- 7 **end**

computes  
the update

computes the  
residual

## DUMBO: NYU's YARN Cluster

- 48 total nodes, running Cloudera CDH 5.9.0 (Hadoop 2.6.0 with YARN)
- Two of the nodes are the Master nodes (babar and hathi).
- Two other nodes are the login nodes (dumbo0 and dumbo1)
- That leaves us with 44 compute nodes.
- Each node has 128 GB RAM.
- Each node has two CPUs. Each CPU is a 8-core Intel 'Haswell' (c 2014).

## Spark+YARN environment settings

- **'-executor-cores'** : Tells Spark to assign each executor 5 cores. Each executor is a JVM instance and will run 5 concurrent tasks.

## Spark+YARN environment settings

- **'-executor-cores'** : Tells Spark to assign each executor 5 cores. Each executor is a JVM instance and will run 5 concurrent tasks.
- **'-num-executors'** : This is the number of nodes in abstract.  $15/5 = 3$  so  $3 \times 43 = 129$  executors.

## Spark+YARN environment settings

- **'-executor-cores'** : Tells Spark to assign each executor 5 cores. Each executor is a JVM instance and will run 5 concurrent tasks.
- **'-num-executors'** : This is the number of nodes in abstract.  $15/5 = 3$  so  $3 \times 43 = 129$  executors.
- **'-executor-memory'**: We have 127 GB RAM and 3 executors per node  $\rightarrow$  each executor gets  $\sim 42$  GB of RAM. Accounting for  $\sim 10\%$  overhead memory  $\rightarrow$  each executor get 38 GB of RAM .

## First implementation: Using Spark's Distributed

- Use Spark's (MLlib) distributed Matrix, specifically BlockMatrix and CoordinateMatrix to implement  $A$  and  $u$  as dense matrices.

## Second implementation : Using scalable matrix multiplication

$$\begin{bmatrix} 1 & 1 & -1 \\ 4 & 0 & 2 \\ 1 & 0 & 0 \end{bmatrix} \cdot \begin{bmatrix} 2 & -1 \\ 3 & -2 \\ 0 & 1 \end{bmatrix} = \begin{bmatrix} 5 & -4 \\ 8 & -2 \\ 2 & -1 \end{bmatrix}$$

**A**                      **B**                      **C**

Co-group columns of  $A$  with rows of  $B$  and perform a outer product on each group. Reduce by row to compute  $C$ :

$$\begin{array}{|c|} \hline \begin{bmatrix} 1 \\ 4 \\ 1 \end{bmatrix}, \begin{bmatrix} 2 & -1 \end{bmatrix} \\ \hline \begin{bmatrix} 1 \\ 0 \\ 0 \end{bmatrix}, \begin{bmatrix} 3 & -2 \end{bmatrix} \\ \hline \begin{bmatrix} -1 \\ 2 \\ 0 \end{bmatrix}, \begin{bmatrix} 0 & 1 \end{bmatrix} \\ \hline \end{array} \cdot \text{map}(\_ \otimes \_) = \begin{array}{|c|} \hline \begin{bmatrix} 2 & -1 \\ 8 & -4 \\ 2 & -1 \end{bmatrix} \\ \hline \begin{bmatrix} 3 & -2 \\ 0 & 0 \\ 0 & 0 \end{bmatrix} \\ \hline \begin{bmatrix} 0 & -1 \\ 0 & 2 \\ 0 & 0 \end{bmatrix} \\ \hline \end{array} \cdot \text{reduce}(\_ + \_) = \begin{bmatrix} 5 & -4 \\ 8 & -2 \\ 2 & -1 \end{bmatrix}$$

## Second implementation

Computing  $A \cdot u_k$

```
val Auk_rdd = A_rdd.join(u_rdd)
    .flatMap{case(k, v) => v._1.map(mv => (mv._1, mv._2 * v._2))}
    .reduceByKey(_ + _)
```



## Third implementation: Don't move $A$ , broadcast $u$

```
var u_recie = sc.broadcast(u_rdd.collect().toMap).value
```

Now, each core has a complete copy of  $u$ . Compute  $Au$  with one less shuffling:

```
Auk_rdd = A_rdd
  mapPartitions({ iter =>
    for {(ak, av) <- iter}
      yield (ak, (av, u_recie.get(ak).get))
  }, preservesPartitioning = true)
  .flatMap{ case (k, v) => v._1.map(Av => (Av._1, Av._2 * v._2)) }
  .reduceByKey(_ + _).cache() // reduce on rows
```

## Third implementation: more improvements

- Avoid redundant computation of  $Au^{k+1}$  : Better solution: compute  $Au^{k+1}$  in  $k$ th iteration and cache it for the next iteration.

## Third implementation: more improvements

- Avoid redundant computation of  $Au^{k+1}$  : Better solution: compute  $Au^{k+1}$  in  $k$ th iteration and cache it for the next iteration.
- Use accumulator to compute the residual: Accumulators are added to through an associative and commutative operation and therefore are efficient.

## Third implementation: more improvements

- Avoid redundant computation of  $Au^{k+1}$  : Better solution: compute  $Au^{k+1}$  in  $k$ th iteration and cache it for the next iteration.
- Use accumulator to compute the residual: Accumulators are added to through an associative and commutative operation and therefore are efficient.
- Increase 'numTasks' for reduceByKey operations (default is 8).

## Third implementation: every iteration:

```
49 //      Computing  $u_{k+1} = 0.5 \cdot hsq - 0.5 \cdot u_k + u_k$ 
50       $u\_rdd = Auk\_rdd.mapValues(hsqHalf - 0.5 * \_)$ 
51      .mapPartitions({ iter =>
52          for {(ak, av) <- iter
53              } yield (ak, (av, u_recie.get(ak).get))
54          }, preservesPartitioning = true)
55      .mapValues{ s => s._1 + s._2 }.cache()
56       $u\_recie = sc.broadcast(u\_rdd.collect().toMap).value$ 
57 //      Computing  $A \cdot u_{k+1}$ 
58       $Auk\_rdd = A\_rdd$ 
59      .mapPartitions({ iter =>
60          for {(ak, av) <- iter
61              } yield (ak, (av, u_recie.get(ak).get))
62          }, preservesPartitioning = true)
63      .flatMap{ case (k, v) => v._1.map(Av => (Av._1, Av._2 * v._2)) }
64      .reduceByKey(_ + _, numTasks).cache() // reduce on rows
65
66 //      Computing residual
67       $val accum = sc.doubleAccumulator("Accumulating residual")$ 
68       $Auk\_rdd.mapValues(invhsq * \_ - 1)$ 
69      .map{ case (k, v) => v * v }
70      .foreach(x => accum.add(x))
71       $ress(k) = Math.sqrt(accum.value)$ 
72       $k += 1$ 
```

update u: without shuffling.

broadcast u

Compute new Au and cache it

use accumulator for residual

## Third implementation: Spark UI view

- Every SparkContext launches a web UI, by default on port 4040.
- Spark can only run one concurrent task for every partition of an RDD, up to the number of cores (either physical or spark cores).
- Running previous code for  $N = 50000$ , 1 iteration and 5 partitions on my laptop : `master("local[*"])`

### Completed Jobs (3)

Job Id	Description	Submitted	Duration	Stages: Succeeded/Total	Tasks (for all stages): Succeeded/Total
2	<a href="#">foreach at Jacobi.scala:70</a>	2017/05/15 19:25:49	0.4 s	2/2 (1 skipped)	10/10 (5 skipped)
1	<a href="#">collect at Jacobi.scala:56</a>	2017/05/15 19:25:47	2 s	3/3	15/15
0	<a href="#">collect at Jacobi.scala:38</a>	2017/05/15 19:25:45	0.5 s	1/1	8/8

# Third implementation: Spark UI view


## Details for Job 1

Status: SUCCEEDED

Completed Stages: 3

▶ Event Timeline

▼ DAG Visualization



## Completed Stages (3)

Stage Id	Description		Submitted	Duration	Tasks: Succeeded/Total	Input	Output	Shuffle Read	Shuffle Write
3	<a href="#">collect at Jacobi.scala:56</a>	<a href="#">+details</a>	2017/05/15 19:25:49	0.4 s	5/5			599.7 KB	
2	<a href="#">flatMap at Jacobi.scala:45</a>	<a href="#">+details</a>	2017/05/15 19:25:47	1 s	5/5			1236.3 KB	599.7 KB
1	<a href="#">parallelize at Jacobi.scala:32</a>	<a href="#">+details</a>	2017/05/15 19:25:47	0.4 s	5/5				1236.3 KB